

Yuwen Sang

A8: Mongo

Part1: Queries

Q2: pid, state, and name of the children for people who have 7 children. Do not include object ids.

```
>db.thePeople.aggregate([{$match:{numChildren:{$eq:7}}},{ $project:{_id:0, firstName:0, MI:0, lastName:0, age:0, birth:0, salary:0, numChildren:0, weight:0, height:0}}])
```

```
[ { pid: 3,
  state: 'TX',
  children:
    [ 'Victoria M Brown',
      'John I Brown',
      'Priya J Brown',
      'Rahul G Brown',
      'Michelle G Brown',
      'Linda Y Brown',
      'Jennifer B Brown' ] },
  { pid: 513,
    state: 'OR',
    children:
      [ 'Michelle W Jackson',
        'David A Jackson',
        'Victoria M Jackson',
        'Santiago P Jackson',
        'Isabella H Jackson',
        'Mohammed G Jackson',
        'Amy I Jackson' ] },
  { pid: 612,
    state: 'SD',
    children:
      [ 'Joseph Q Wilson',
        'Julie A Wilson',
        'David X Wilson',
        'Paul A Wilson',
        'Sofia O Wilson',
        'Amy Z Wilson',
        'Amy Z Wilson' ] },
  { pid: 745,
    state: 'IN',
    children:
      [ 'Amy T Zhao',
        'Victoria E Zhao',
        'Jayden E Zhao',
        'Kayla V Zhao',
        'Madison E Zhao',
        'Jayden I Zhao',
        'Neha D Zhao' ] },
  { pid: 816,
    state: 'MD',
    children:
      [ 'Demarco Q Le',
        'Vicky T Le',
        'Isabella G Le',
        'Jennifer R Le',
```

```
'Sofia Y Le',
'John Y Le',
'Diego H Le' ]}],
{ pid: 881,
  state: 'NE',
  children:
    [ 'Madison U Jones',
      'Priya F Jones',
      'Amy L Jones',
      'Kayla T Jones',
      'John F Jones',
      'Isabella X Jones',
      'Paul H Jones' ] }],
{ pid: 993,
  state: 'OR',
  children:
    [ 'Ashley C Wu',
      'Demarco R Wu',
      'Hasan C Wu',
      'David N Wu',
      'Ava U Wu',
      'Jeffrey V Wu',
      'Neha H Wu' ] }],
{ pid: 1010,
  state: 'MS',
  children:
    [ 'Bob X Park',
      'Mohammad I Park',
      'Amy F Park',
      'Ashley N Park',
      'Santiago Z Park',
      'Julie D Park',
      'Amy O Park' ] }],
{ pid: 1240,
  state: 'CT',
  children:
    [ 'Jayden V Wong',
      'Demarco U Wong',
      'Julie R Wong',
      'Julie F Wong',
      'Hannah K Wong',
      'Jeffrey Z Wong',
      'Julie O Wong' ] }],
{ pid: 1321,
  state: 'KS',
  children:
    [ 'Hasan V Anderson',
      'Madison Z Anderson',
      'Kayla M Anderson',
      'Jayden D Anderson',
      'John D Anderson',
      'Julia J Anderson',
      'John N Anderson' ] }],
{ pid: 1343,
  state: 'KY',
  children:
    [ 'Jennifer C Tanaka',
      'Amy R Tanaka',
```

```

    'Hannah U Tanaka',
    'Isabella I Tanaka',
    'Michelle Z Tanaka',
    'Mohammed D Tanaka',
    'Mohammad P Tanaka' ] },
{ pid: 1537,
  state: 'KS',
  children:
    [ 'John Z Sato',
      'Madison F Sato',
      'Jeffrey T Sato',
      'Demarco P Sato',
      'Santiago I Sato',
      'Victoria A Sato',
      'Jennifer K Sato' ] },
{ pid: 1791,
  state: 'VT',
  children:
    [ 'Julie K Williams',
      'David Z Williams',
      'Vivek G Williams',
      'Madison B Williams',
      'Rahul N Williams',
      'Alejandro A Williams',
      'Alejandro R Williams' ] },
{ pid: 1879,
  state: 'WV',
  children:
    [ 'Noah Y Jones',
      'Paul K Jones',
      'Noah B Jones',
      'Michelle B Jones',
      'Sofia Y Jones',
      'Diego V Jones',
      'Victoria R Jones' ] },
{ pid: 1994,
  state: 'AL',
  children:
    [ 'Priya D Lee',
      'Mary D Lee',
      'Kayla G Lee',
      'Isabella O Lee',
      'Priya H Lee',
      'Ashley H Lee',
      'Jayden Z Lee' ] ] ]

```

Q4: Complete info of people who live in CA and have 5 or 6 children

```
>db.thePeople.aggregate([{$match:{$or:[{numChildren:{$eq:5}},{numChildren:{$eq:6}}]}},{ $match:{state:{$eq:"CA"}}}])
```

```

[ { _id: ObjectId("604fe1a25701ce73669814c4"),
  pid: 169,
  firstName: 'Amit',
  MI: 'X',
  lastName: 'Lee',
  state: 'CA',

```

```
age: 117,  
birth: 1902,  
salary: 61614,  
numChildren: 6,  
children:  
[ 'Sarah G Lee',  
  'Madison D Lee',  
  'Victoria O Lee',  
  'Isabella C Lee',  
  'William O Lee',  
  'Ava Q Lee' ],  
weight: 114,  
height: 180 },  
{ _id: ObjectId("604fe1a35701ce7366981701"),  
  pid: 742,  
  firstName: 'Demarco',  
  MI: 'Y',  
  lastName: 'Jackson',  
  state: 'CA',  
  age: 84,  
  birth: 1935,  
  salary: 67539,  
  numChildren: 5,  
  children:  
  [ 'Bob F Jackson',  
    'Daniella I Jackson',  
    'Peter U Jackson',  
    'Julie F Jackson',  
    'Neha J Jackson' ],  
  weight: 134,  
  height: 191 },  
{ _id: ObjectId("604fe1a35701ce73669817dc"),  
  pid: 961,  
  firstName: 'Fatimah',  
  MI: 'X',  
  lastName: 'Baker',  
  state: 'CA',  
  age: 84,  
  birth: 1935,  
  salary: 118150,  
  numChildren: 5,  
  children:  
  [ 'Victoria V Baker',  
    'Mary B Baker',  
    'Bob F Baker',  
    'Kayla L Baker',  
    'Amit R Baker' ],  
  weight: 62,  
  height: 197 },  
{ _id: ObjectId("604fe1a45701ce736698199b"),  
  pid: 1408,  
  firstName: 'Mary',  
  MI: 'X',  
  lastName: 'Brown',  
  state: 'CA',  
  age: 30,  
  birth: 1989,  
  salary: 43826,
```

```

numChildren: 5,
children:
[ 'Hannah Y Brown',
  'Hasan Q Brown',
  'Peter D Brown',
  'Jayla D Brown',
  'William X Brown' ],
weight: 82,
height: 189 } ]

```

Q5: List the pid and children names for all people who have a child whose name contains 'Bob A':

```

>db.thePeople.aggregate([{$match:{$children:{$regex: "Bob A"}}},{ $project: {_id:0, firstName:0, MI:0,
lastName:0, state:0, age:0, birth:0, salary:0, numChildren:0, weight:0, height:0}}])

```

```

[ { pid: 692,
  children:
  [ 'Bob A Chen',
    'Linda P Chen',
    'Madison Z Chen',
    'Linda P Chen',
    'Jennifer D Chen' ] },
  { pid: 991,
    children:
    [ 'Mohammad N Chan',
      'Bob A Chan',
      'Madison T Chan',
      'Victoria N Chan' ] },
  { pid: 1612,
    children:
    [ 'Bob A Takahashi',
      'Neha H Takahashi',
      'Jennifer L Takahashi',
      'Hannah C Takahashi' ] } ]

```

Q9: Aggregation: average/min/max salary for midwest state, where I am assuming there are 12 midwest states:

```

>db.thePeople.aggregate([{$match:{$or: [{state:{$eq:"ND"}},{state:{$eq:"SD"}},{state:{$eq:"NE"}},{state:{$eq:"KS"}},{state:{$eq:"MN"}},{state:{$eq:"IA"}},{state:{$eq:"MS"}},{state:{$eq:"WI"}},{state:{$eq:"IL"}},{state:{$eq:"IN"}},{state:{$eq:"MI"}},{state:{$eq:"OH"}}]}},{ $group: {_id:"$state", avgSalary:{$avg:"$salary"}, minSalary:{$min:"$salary"}, maxSalary:{$max:"$salary"}, "numInGroup":{$sum:1}}},{ $sort: {"_id":1}}])

```

```

[ { _id: 'IA',
  avgSalary: 83429.84615384616,
  minSalary: 32606,
  maxSalary: 128314,
  numInGroup: 39 },
  { _id: 'IL',
    avgSalary: 74737.85294117648,
    minSalary: 31481,
    maxSalary: 114905,
    numInGroup: 34 },

```

```

{ _id: 'IN',
  avgSalary: 85336.21212121213,
  minSalary: 36705,
  maxSalary: 129445,
  numInGroup: 33 },
{ _id: 'KS',
  avgSalary: 76605.38888888889,
  minSalary: 37380,
  maxSalary: 125949,
  numInGroup: 36 },
{ _id: 'MI',
  avgSalary: 78596.93023255814,
  minSalary: 30837,
  maxSalary: 128545,
  numInGroup: 43 },
{ _id: 'MN',
  avgSalary: 79549.45714285714,
  minSalary: 33945,
  maxSalary: 126530,
  numInGroup: 35 },
{ _id: 'MS',
  avgSalary: 82858.68,
  minSalary: 30591,
  maxSalary: 124771,
  numInGroup: 25 },
{ _id: 'ND',
  avgSalary: 84828.57142857143,
  minSalary: 31346,
  maxSalary: 126747,
  numInGroup: 42 },
{ _id: 'NE',
  avgSalary: 76561.78723404255,
  minSalary: 31571,
  maxSalary: 125839,
  numInGroup: 47 },
{ _id: 'OH',
  avgSalary: 81283.39130434782,
  minSalary: 32641,
  maxSalary: 129161,
  numInGroup: 46 },
{ _id: 'SD',
  avgSalary: 81643.41666666667,
  minSalary: 35200,
  maxSalary: 128021,
  numInGroup: 36 },
{ _id: 'WI',
  avgSalary: 79170.7435897436,
  minSalary: 30949,
  maxSalary: 126922,
  numInGroup: 39 } ]

```

Q10: Aggregation: average salary in states where the average salary within that state is $\geq 82,000$ and how many people in the grouping for each state:

```

>db.thePeople.aggregate([{$group:{_id:"$state",avgSalary:{$avg:"$salary"},numInGroup:{$sum:1}}},{match:{avgSalary:{$gte:82000}}},{sort:{ "_id":1}}])

```

```
[ { _id: 'CO', avgSalary: 82695.52380952382, numInGroup: 42 },
  { _id: 'IA', avgSalary: 83429.84615384616, numInGroup: 39 },
  { _id: 'IN', avgSalary: 85336.21212121213, numInGroup: 33 },
  { _id: 'MD', avgSalary: 87204.64406779662, numInGroup: 59 },
  { _id: 'MS', avgSalary: 82858.68, numInGroup: 25 },
  { _id: 'MT', avgSalary: 83048.78048780488, numInGroup: 41 },
  { _id: 'NC', avgSalary: 83518.30555555556, numInGroup: 36 },
  { _id: 'ND', avgSalary: 84828.57142857143, numInGroup: 42 },
  { _id: 'NH', avgSalary: 83915.575, numInGroup: 40 },
  { _id: 'NM', avgSalary: 83036.27777777778, numInGroup: 36 },
  { _id: 'NY', avgSalary: 82193.19565217392, numInGroup: 46 },
  { _id: 'OK', avgSalary: 85440.54838709677, numInGroup: 31 },
  { _id: 'PA', avgSalary: 85585.52173913043, numInGroup: 46 },
  { _id: 'SC', avgSalary: 85474.84615384616, numInGroup: 39 },
  { _id: 'VT', avgSalary: 89343.45714285714, numInGroup: 35 },
  { _id: 'WV', avgSalary: 87893.84210526316, numInGroup: 38 } ]
```

Q11: Aggregation: average/min/max salary for midwest state whose average salary > 82000:

```
>db.thePeople.aggregate([{$match:{$or:[{state:{$eq:"ND"}},{state:{$eq:"SD"}},{state:{$eq:"NE"}},{state:{$eq:"KS"}},{state:{$eq:"MN"}},{state:{$eq:"IA"}},{state:{$eq:"MS"}},{state:{$eq:"WI"}},{state:{$eq:"IL"}},{state:{$eq:"IN"}},{state:{$eq:"MI"}},{state:{$eq:"OH"}}]}},
{$group:{_id:"$state",avgSalary:{$avg:"$salary"},minSalary:{$min:"$salary"},maxSalary:{$max:"$salary"},numInGroup:{$sum:1}}}, {$match:{avgSalary:{$gt:82000}}},{$sort:{"_id":1}}])
```

```
[ { _id: 'IA',
  avgSalary: 83429.84615384616,
  minSalary: 32606,
  maxSalary: 128314,
  numInGroup: 39 },
  { _id: 'IN',
  avgSalary: 85336.21212121213,
  minSalary: 36705,
  maxSalary: 129445,
  numInGroup: 33 },
  { _id: 'MS',
  avgSalary: 82858.68,
  minSalary: 30591,
  maxSalary: 124771,
  numInGroup: 25 },
  { _id: 'ND',
  avgSalary: 84828.57142857143,
  minSalary: 31346,
  maxSalary: 126747,
  numInGroup: 42 } ]
```

Part2: Updates/Deletes

1. Change 1 Tuple:

Based on the result of the query 2, change the people whose pid = 3, its state to "CA".

Before Update, the people whose id = 3, was from "TX" (see part1 query 2).

Then, after input the command in MongoDB:

```
>db.thePeople.updateOne({pid:3},{set:{state:"CA"}})
```

Run Q2 command again, we get, for pid= 3:

```
[ { pid: 3,
  state: 'CA',
  children:
    [ 'Victoria M Brown',
      'John I Brown',
      'Priya J Brown',
      'Rahul G Brown',
      'Michelle G Brown',
      'Linda Y Brown',
      'Jennifer B Brown' ] },
```

2. Change Multiple Tuples

Let's change the people living in CA their first name to "ABC".

```
>db.thePeople.updateMany({state:{$eq:"CA"}},{ $set:{firstName: "ABC"}})
```

After the update, we get the new result of Q4 will be:

```
[ { _id: ObjectId("604fela25701ce73669814c4"),
  pid: 169,
  firstName: 'ABC',
  MI: 'X',
  lastName: 'Lee',
  state: 'CA',
  age: 117,
  birth: 1902,
  salary: 61614,
  numChildren: 6,
  children:
    [ 'Sarah G Lee',
      'Madison D Lee',
      'Victoria O Lee',
      'Isabella C Lee',
      'William O Lee',
      'Ava Q Lee' ],
  weight: 114,
  height: 180 },
  { _id: ObjectId("604fela35701ce7366981701"),
  pid: 742,
  firstName: 'ABC',
  MI: 'Y',
  lastName: 'Jackson',
  state: 'CA',
  age: 84,
  birth: 1935,
  salary: 67539,
  numChildren: 5,
  children:
    [ 'Bob F Jackson',
      'Daniella I Jackson',
      'Peter U Jackson',
      'Julie F Jackson',
      'Neha J Jackson' ],
  weight: 134,
  height: 191 },
  { _id: ObjectId("604fela35701ce73669817dc"),
  pid: 961,
  firstName: 'ABC',
  MI: 'X',
```



```

    lastName: 'Baker',
    state: 'CA',
    age: 84,
    birth: 1935,
    salary: 118150,
    numChildren: 5,
    children:
      [ 'Victoria V Baker',
        'Mary B Baker',
        'Bob F Baker',
        'Kayla L Baker',
        'Amit R Baker' ],
    weight: 62,
    height: 197 },
{ _id: ObjectId("604fe1a45701ce736698199b"),
  pid: 1408,
  firstName: 'ABC',
  MI: 'X',
  lastName: 'Brown',
  state: 'CA',
  age: 30,
  birth: 1989,
  salary: 43826,
  numChildren: 5,
  children:
    [ 'Hannah Y Brown',
      'Hasan Q Brown',
      'Peter D Brown',
      'Jayla D Brown',
      'William X Brown' ],
  weight: 82,
  height: 189 } ]

```

3. Delete multiple documents

For query 2, we can see there are 2 people living in KS and having 7 children. Let's delete them:

```
> db.thePeople.deleteMany({$and:[{ numChildren:{ $eq:7 }},{state:{ $eq:"KS" } } ]})
```

```
{ acknowledged: true, deletedCount: 2 }
```

Then run Q2 again:

```

[ { pid: 3,
  state: 'CA',
  children:
    [ 'Victoria M Brown',
      'John I Brown',
      'Priya J Brown',
      'Rahul G Brown',
      'Michelle G Brown',
      'Linda Y Brown',
      'Jennifer B Brown' ] },
  { pid: 513,
    state: 'OR',
    children:
      [ 'Michelle W Jackson',
        'David A Jackson',
        'Victoria M Jackson',
        'Santiago P Jackson',
        'Isabella H Jackson',
        'Mohammed G Jackson',
        'Amy I Jackson' ] },

```

```
{ pid: 612,
  state: 'SD',
  children:
    [ 'Joseph Q Wilson',
      'Julie A Wilson',
      'David X Wilson',
      'Paul A Wilson',
      'Sofia O Wilson',
      'Amy Z Wilson',
      'Amy Z Wilson' ] },
{ pid: 745,
  state: 'IN',
  children:
    [ 'Amy T Zhao',
      'Victoria E Zhao',
      'Jayden E Zhao',
      'Kayla V Zhao',
      'Madison E Zhao',
      'Jayden I Zhao',
      'Neha D Zhao' ] },
{ pid: 816,
  state: 'MD',
  children:
    [ 'Demarco Q Le',
      'Vicky T Le',
      'Isabella G Le',
      'Jennifer R Le',
      'Sofia Y Le',
      'John Y Le',
      'Diego H Le' ] },
{ pid: 881,
  state: 'NE',
  children:
    [ 'Madison U Jones',
      'Priya F Jones',
      'Amy L Jones',
      'Kayla T Jones',
      'John F Jones',
      'Isabella X Jones',
      'Paul H Jones' ] },
{ pid: 993,
  state: 'OR',
  children:
    [ 'Ashley C Wu',
      'Demarco R Wu',
      'Hasan C Wu',
      'David N Wu',
      'Ava U Wu',
      'Jeffrey V Wu',
      'Neha H Wu' ] },
{ pid: 1010,
  state: 'MS',
  children:
    [ 'Bob X Park',
      'Mohammad I Park',
      'Amy F Park',
      'Ashley N Park',
      'Santiago Z Park',
      'Julie D Park',
      'Amy O Park' ] },
{ pid: 1240,
  state: 'CT',
  children:
```

```

    [ 'Jayden V Wong',
      'Demarco U Wong',
      'Julie R Wong',
      'Julie F Wong',
      'Hannah K Wong',
      'Jeffrey Z Wong',
      'Julie O Wong' ] },
  { pid: 1343,
    state: 'KY',
    children:
      [ 'Jennifer C Tanaka',
        'Amy R Tanaka',
        'Hannah U Tanaka',
        'Isabella I Tanaka',
        'Michelle Z Tanaka',
        'Mohammed D Tanaka',
        'Mohammad P Tanaka' ] },
  { pid: 1791,
    state: 'VT',
    children:
      [ 'Julie K Williams',
        'David Z Williams',
        'Vivek G Williams',
        'Madison B Williams',
        'Rahul N Williams',
        'Alejandro A Williams',
        'Alejandro R Williams' ] },
  { pid: 1879,
    state: 'WV',
    children:
      [ 'Noah Y Jones',
        'Paul K Jones',
        'Noah B Jones',
        'Michelle B Jones',
        'Sofia Y Jones',
        'Diego V Jones',
        'Victoria R Jones' ] },
  { pid: 1994,
    state: 'AL',
    children:
      [ 'Priya D Lee',
        'Mary D Lee',
        'Kayla G Lee',
        'Isabella O Lee',
        'Priya H Lee',
        'Ashley H Lee',
        'Jayden Z Lee' ] } ] ]

```

Part 3: Indexing

Now let's expand the database into 200,000, and find the people that has more than 1 children.

Then, before using index, the query is:

```
>db.thePeople.find({numChildren:{>1}}).explain("executionStats")
```

The result is:

```
{ queryPlanner:
  { plannerVersion: 1,
```

```

namespace: 'db_people.thePeople',
indexFilterSet: false,
parsedQuery: { numChildren: { '$gt': 1 } },
winningPlan:
  { stage: 'COLLSCAN',
    filter: { numChildren: { '$gt': 1 } },
    direction: 'forward' },
rejectedPlans: [] },
executionStats:
  { executionSuccess: true,
    nReturned: 173181,
    executionTimeMillis: 348,
    totalKeysExamined: 0,
    totalDocsExamined: 200000,
    executionStages:
      { stage: 'COLLSCAN',
        filter: { numChildren: { '$gt': 1 } },
        nReturned: 173181,
        executionTimeMillisEstimate: 26,
        works: 200002,
        advanced: 173181,
        needTime: 26820,
        needYield: 0,
        saveState: 200,
        restoreState: 200,
        isEOF: 1,
        direction: 'forward',
        docsExamined: 200000 } },
serverInfo:
  { host: 'LAPTOP-ANQ85J7N',
    port: 27017,
    version: '4.4.4',
    gitVersion: '8db30a63db1a9d84bdcad0c83369623f708e0397' },
ok: 1 }

```

,which spend 348 milsec.

Now create an index:

```
>db.thePeople.ensureIndex({numChildren:1})
```

We get return:

```

< { createdCollectionAutomatically: false,
  numIndexesBefore: 1,
  numIndexesAfter: 2,
  ok: 1 }

```

Test the index that we just created:

```

> db.thePeople.getIndexes()
< [ { v: 2, key: { _id: 1 }, name: '_id_ ' },
    { v: 2, key: { numChildren: 1 }, name: 'numChildren_1' } ]

```

Run the query again:

```
>db.thePeople.find({numChildren:{$gt:1}}).explain("executionStats")
```

We get the result:

```
{ queryPlanner:
  { plannerVersion: 1,
    namespace: 'db_people.thePeople',
    indexFilterSet: false,
    parsedQuery: { numChildren: { '$gt': 1 } },
    winningPlan:
      { stage: 'FETCH',
        inputStage:
          { stage: 'IXSCAN',
            keyPattern: { numChildren: 1 },
            indexName: 'numChildren_1',
            isMultiKey: false,
            multiKeyPaths: { numChildren: [ ] },
            isUnique: false,
            isSparse: false,
            isPartial: false,
            indexVersion: 2,
            direction: 'forward',
            indexBounds: { numChildren: [ '(1, inf.0]' ] } },
        rejectedPlans: [ ] },
    executionStats:
      { executionSuccess: true,
        nReturned: 173181,
        executionTimeMillis: 27,
        totalKeysExamined: 173181,
        totalDocsExamined: 173181,
        executionStages:
          { stage: 'FETCH',
            nReturned: 173181,
            executionTimeMillisEstimate: 7,
            works: 173182,
            advanced: 173181,
            needTime: 0,
            needYield: 0,
            saveState: 173,
            restoreState: 173,
            isEOF: 1,
            docsExamined: 173181,
            alreadyHasObj: 0,
            inputStage:
              { stage: 'IXSCAN',
                nReturned: 173181,
                executionTimeMillisEstimate: 3,
                works: 173182,
                advanced: 173181,
                needTime: 0,
                needYield: 0,
                saveState: 173,
                restoreState: 173,
                isEOF: 1,
                keyPattern: { numChildren: 1 },
                indexName: 'numChildren_1',
                isMultiKey: false,
                multiKeyPaths: { numChildren: [ ] },
                isUnique: false,
                isSparse: false,
                isPartial: false,
                indexVersion: 2,
```

```
    direction: 'forward',
    indexBounds: { numChildren: [ '(1, inf.0]' ] },
    keysExamined: 173181,
    seeks: 1,
    dupsTested: 0,
    dupsDropped: 0 } } },
serverInfo:
  { host: 'LAPTOP-ANQ85J7N',
    port: 27017,
    version: '4.4.4',
    gitVersion: '8db30a63db1a9d84bdcad0c83369623f708e0397' },
ok: 1 }
```

, which using indexing, the *execution time* reduced from 348 to 27, *execution time millis* estimate from 26 to 7 milsec.